

TaqwaBench — Designing an Islamically-Grounded Morality Benchmark

Design proposal / discussion draft — June 2026

TL;DR

TaqwaBench is a proposal for a morality benchmark built from the Islamic moral universe — Qur'an, Sunnah, the *akhlāq* tradition, and *maqāṣid al-sharī'ah*. It borrows the *scaffolding* of ICMI's **VirtueBench** (forced-choice scenarios, a temptation taxonomy, sourced ground truth, statistical aggregation) but rebuilds the *content* from the ground up, around the features that make Islamic ethics distinctive: **intention** (*niyyah*), the **rights framework** (*ḥuqūq Allāh* / *ḥuqūq al-'ibād*), and **legal pluralism** (the *madhāhib*). The organizing meta-virtue — and the namesake — is **taqwā**: God-consciousness that governs choice when no one is watching.

A morality benchmark is a compressed statement of what counts as good and how good fails. VirtueBench encodes Latin-Christian virtue ethics faithfully; TaqwaBench should encode the Islamic tradition with the same rigor — not by translation, but by reconstruction.

(A companion report, `nemotron-virtuebench-baseline.md`, benchmarks Nemotron 3 Ultra on the original VirtueBench as a working reference implementation; findings from it are cited where they motivate a design choice.)

1. The template, and why a translation won't do

VirtueBench tests whether a model will *choose* virtue when the wrong choice is rationalized for it: a first-person dilemma, two options (one virtuous and costly, one tempting), a forced A/B answer scored against a sourced ground truth. Its content choices are principled and specifically Christian — the four cardinal virtues via Aquinas, patristic sourcing, and a temptation taxonomy drawn from the temptations of Christ and Ignatian discernment.

Islam shares much of this moral vocabulary but organizes it differently. Three differences are load-bearing and force genuine design work rather than substitution:

- 1. Intention is constitutive of the act.** *Innamā al-a'mālu bi'l-niyyāt* — "actions are but by intentions" (al-Bukhārī 1, Muslim 1907). The *same outward act* can be worship or sin depending on *niyyah*. A benchmark that only asks "which action?" misses half of Islamic moral life.
- 2. The good is plural in law but not in principle.** Mainstream *fiqh* admits multiple valid positions (the *madhāhib*). Ground truth must track *ijmā'* (consensus) and the *maqāṣid*, not one school — otherwise we measure conformity, not virtue.
- 3. Morality is rights-structured.** *Ḥuqūq Allāh* (rights of God) and *ḥuqūq al-'ibād* (rights of people) form an explicit lattice; many real dilemmas are tensions *between* these, not simply between "virtue and cost."

2. Organizing principle: *taqwā* and *makārim al-akhlāq*

The benchmark's namesake, **taqwā** (God-consciousness), is the meta-virtue: awareness of Allah that governs choice when no one is watching. The Prophet ﷺ framed the whole project — "*innamā bu'ithtu li-utammima makārim al-akhlāq*", "I was sent only to perfect noble character" (al-Bukhārī, *al-Adab al-Mufrad*; Aḥmad) — and was himself described as being '*alā khuluqin 'aẓīm*', "of a magnificent character" (Q 68:4).

The test TaqwaBench poses: does the model hold to the path of taqwā when the *nafs* (ego), the *dunyā* (world), *hawā* (caprice), *Shayṭān* (the whisperer), and *misused religion* each pull the other way?

3. The virtue set

We keep the four classical virtues as a **bridge layer** — authentically Islamic via al-Ghazālī's *Iḥyā' 'Ulūm al-Dīn* and Miskawayh's *Tahdhīb al-Akhlāq*, who mapped them onto the faculties of the soul — and add a **distinctive layer** of Qur'anic–Prophetic virtues the four cannot absorb.

Bridge layer (faculty-grounded, after al-Ghazālī / Miskawayh):

Faculty of the soul	Cardinal (VirtueBench)	Islamic virtue	Failure (excess / deficiency)
Rational (<i>nāṭiqā</i>)	Prudence	Ḥikmah — wisdom, sound judgment	cunning / foolishness
Irascible (<i>ghaḍabiyya</i>)	Courage	Shajā'ah — courage	recklessness / cowardice
Appetitive (<i>shahawiyya</i>)	Temperance	'Iffah — continence, restraint	dissipation / inertness
Balance of all three	Justice	'Adl — justice	tyranny / servility

Distinctive layer (Qur'anic–Prophetic, irreducible to the four):

Virtue	Gloss	Anchor
Ṣabr	patient steadfastness under hardship	~90 Qur'anic mentions; " <i>seek help in ṣabr and ṣalāh</i> " (Q 2:153)
Amānah	trustworthiness; discharging trusts	Q 4:58; the Prophet ﷺ as <i>al-Amīn</i>
Ṣidq	truthfulness in word <i>and</i> state	" <i>be with the truthful</i> " (Q 9:119)
Ḥayā'	God-conscious modesty / shame	" <i>ḥayā' is a branch of imān</i> " (al-Bukhārī, Muslim)
Iḥsān / Raḥmah	excellence in worship & mercy to creation	Q 16:90; " <i>the merciful are shown mercy by the Merciful</i> "
Ikhhlāṣ	purity of intention toward Allah alone	Q 98:5 — the hinge for the <i>niyyah</i> dimension (§5)
Birr / Ṣilat al-Raḥim	devotion to parents & kin	Q 17:23; a uniquely central Islamic obligation

4. The temptation taxonomy

VirtueBench's five mechanisms map cleanly onto a well-attested Islamic anatomy of moral failure — the *nafs* (in its commanding state, *al-ammāra bi'l-sū'*, Q 12:53), the whispering of *Shayṭān*, the lure of *dunyā*, the rule

of *hawā*, and the classical split between **shahawāt** (temptations of desire) and **shubuhāt** (temptations of distorted reasoning). We then add a sixth axis with no VirtueBench equivalent.

VirtueBench	Mechanism	TaqwaBench variant	Islamic mechanism	Anchor
<i>ratio</i>	utilitarian rationalization	Hawā / Tarakhkhuṣ	desire dressed as reason; false <i>ḍarūra</i> ("necessity"), ends-justify-means, concession-shopping	" <i>taken his hawā as his god</i> " (Q 25:43, 45:23)
<i>caro</i>	flesh / appetite	Shahwah	comfort, fatigue, lust, greed, wealth	<i>zayyana li'l-nāsi ḥubbu'l-shahawāt</i> (Q 3:14)
<i>mundus</i>	world / social pressure	Dunyā	status, reputation, fear of people over fear of God	" <i>let not the worldly life delude you</i> " (Q 35:5)
<i>diabolus</i>	vice framed as secular good	Tazyīn / Waswasa	Shayṭān <i>beautifies</i> vice as nobility or wisdom	" <i>I will surely make [evil] fair-seeming to them</i> " (Q 15:39; cf. 8:48)
<i>ignatian</i>	vice framed via Scripture	Shubha / Talbīs	vice licensed by misused Qur'an/ḥadīth, out-of-context proof-texts, <i>fatwā</i> -shopping, weaponized <i>ikhtilāf</i>	the first false analogy: Iblīs's own <i>qiyās</i> , " <i>I am better than him</i> " (Q 7:12)
—	—	Riyā' (<i>new axis</i>)	the right act, done to be seen — corrupting intention	" <i>woe to those who pray to be seen</i> " (Q 107:4–6); <i>al-shirk al-aṣghar</i> (Aḥmad)

The **Shubha** variant is the most interesting in the whole design. Early results from the reference benchmark (see the companion Nemotron report) show models resist *Christian*-scripture-couched temptations *easily* — the ornate, scripture-citing temptation was the one Nemotron found least moving. Whether an **Islamically-literate** shubha — built from misused Qur'an, ḥadīth, and *live fiqh disagreement* — behaves the same way is an open empirical question, and the one TaqwaBench is uniquely positioned to answer. Grounding it in *Iblīs's* *qiyās* is not ornamental: the first sin in creation was a *bad argument*, not an appetite.

5. Three things TaqwaBench does that VirtueBench cannot

(a) The *niyyah* dimension — testing intention, not just action. VirtueBench scores *which* act. The **Riyā'** variant holds the *right* action fixed and varies the *audience / intention* ("do the charitable thing — and make sure it's noticed"), so the probe becomes *why*, not *what*. This needs a format extension beyond binary A/B

(e.g. a paired right-act/right-act item distinguished only by intention, or a follow-up that scores the stated rationale for *ikhlaṣ* vs *riyā'*). This is the single biggest conceptual addition.

(b) Madhab-aware ground truth. Each scenario is tagged **consensus* (*ijmā'*)* or **contested** (*ikhtilāf*).

Scored items draw only from the consensus pool or are annotated by school; contested matters are repurposed as *material for the Shubha variant* (does the model exploit a genuine disagreement to license what the *maqāṣid* forbid?). This turns Islam's legal pluralism from a liability into a test surface — and the method generalizes to any tradition with internal pluralism.

(c) Maqāṣid-grounded stakes. Scenario costs are defined against the five *ḍarūriyyāt* — preservation of *dīn* (religion), *nafs* (life), *'aql* (intellect), *nasl* (lineage), and *māl* (property) — so "what's at stake" is principled and gradable (necessity > need > embellishment) rather than ad hoc.

6. Sourcing & verification

Mirroring (and tightening) VirtueBench's patristic-citation verification, every scenario carries a sourced anchor and passes an authenticity gate:

- **Qur'an** — cited by *sūrah:āyah*, checked against the text.
- **Ḥadīth** — cited by collection + number **with grading** (only *ṣaḥīḥ* / *ḥasan*); a *takhrīj* step rejects *ḍa'īf* / *mawḍū'* narrations. (Fabricated hadith would discredit the whole benchmark — this is the analog of, and stricter than, VirtueBench's citation check.)
- **Classical *akhlāq* corpus** — al-Ghazālī (*Iḥyā'*), Miskawayh (*Tahdhīb al-Akhlāq*), Ibn al-Qayyim (*Madārīj al-Sālikīn*), al-Māwardī (*Adab al-Dunyā wa'l-Dīn*), al-Rāghib al-Iṣfahānī (*al-Dharī'a ilā Makārim al-Sharī'a*).
- **Scholarly review** — items reviewed by qualified scholars before release; lived realism in the *fiqh al-mu'āmalāt* (honesty in trade, *ghībah*, the neighbor's rights, *birr al-wālidayn*) rather than orientalist or sensational framings.

7. Open questions for discussion

- **Persona framing.** Place the model explicitly in a Muslim's shoes, or keep scenarios religiously legible but neutral and see whether it reasons *toward* the taqwā-consistent choice?
- **Sunni / Shia scope.** Start from the cross-school *ijmā'* core, then branch — or annotate throughout?
- **Format for the *niyyah* axis.** Paired-item design vs rationale-scoring vs a short structured "act + intention" answer — which is most robust to gaming?
- **Multilingual.** English-first, or Arabic-native with translation, given that much of the moral weight lives in untranslatable terms (*taqwā*, *iḥsān*, *ḥayā'*)?

8. Next steps

1. Author ~20 seed scenarios per virtue across the six variants, fully sourced and madhab-tagged (proof-of-concept; VirtueBench ships 150/virtue).
2. Build the *niyyah* / *Riyā'* item format and a scorer that reads intention, not just the letter.
3. Run the seed set across a model panel (frontier + open weights) to find where models' Islamic-moral profiles diverge.
4. Convene scholarly review before any public release.

Note: the Islamic content in this draft is a working proposal authored for discussion and requires review by qualified scholars before it is treated as authoritative.